

CMPE 249

Report for HW1 – 2D Object Detection

Uttej Kumar Reddy Gade
016065543

Option: Training (1 model) & Inference (2 models)

Models for Training:

YOLOv5 (<https://github.com/ultralytics/yolov5>)

Models for Inference:

FastRCNN (from torchvision.models.detection import fasterrcnn_resnet50_fpn)

YOLOv5 (torch.hub.load('ultralytics/yolov5', 'yolov5s', pretrained=True))

Dataset: [KITTI 2D Object Detection Dataset](#)

Description:

Step 1: Reducing KITTI Dataset for quick training and inference

Extracted the first 1000 images and labels for training and the next 100 for testing from below links.

[Download left color images of object data set \(12 GB\)](#)

[Download training labels of object data set \(5 MB\)](#)

Organised the files in a manner suitable for YOLOv5.

[Directory structure](#)

```
dataset
- images
  -train
  -test
- labels-raw
  -train
  -test
```

TRAINING – YOLOv5

Step 2: Formatting dataset into suitable format for YOLOv5 training

YOLOv5 looks for labels in the format [CATEGORY] [BBOX_X] [BBOX_Y] [HEIGHT] [WIDTH] where BBOX_X and BBOX_Y are the center coordinates of the bounding box. However, the labels downloaded from the source has [CATEGORY] with 14 values encoding more attributes.

So, as the next step, converted the labels to required format using the below link as reference:

<https://github.com/packyan/Kitti2Coco/blob/master/kitti2coco-label-trans.py>

Code:

<https://github.com/uttej कुमारreddy/cmpe249-hw1/blob/master/kitti-labels-to-coco-format.ipynb>

Summary:

Encode all categories in the label to a number and store it.

- a. For each line in label file, extract the 5th-8th value and scale them to get required values.

```
bbox_center_x = float( (x1 + (x2 - x1) / 2.0) / img_width)
bbox_center_y = float( (y1 + (y2 - y1) / 2.0) / img_height)
bbox_width = float((x2 - x1) / img_width)
bbox_height = float((y2 - y1) / img_height)
```

- b. Write them in new label files in the above mentioned format.

Raw Labels:

Pedestrian 0.00 0 -0.20 712.40 143.00 810.73 307.92 1.89 0.48 1.20 1.84 1.47 8.41 0.01

Processed Labels:

0 0.6221936274509804 0.6093513513513513 0.08033496732026148 0.4457297297297298

The processed labels are stored [here](#).

dataset
- labels

Step 3: Downloaded, configured and ran YOLOv5

[Downloaded YOLOv5](#) from Github. Configured the coco.yaml file in data folder as follows:

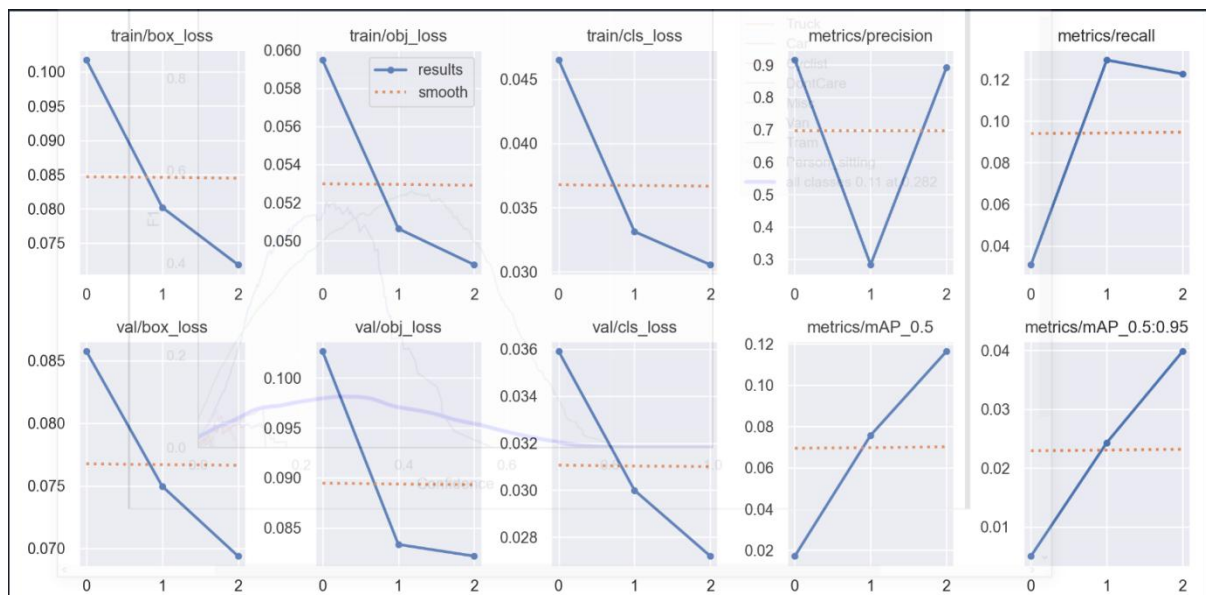
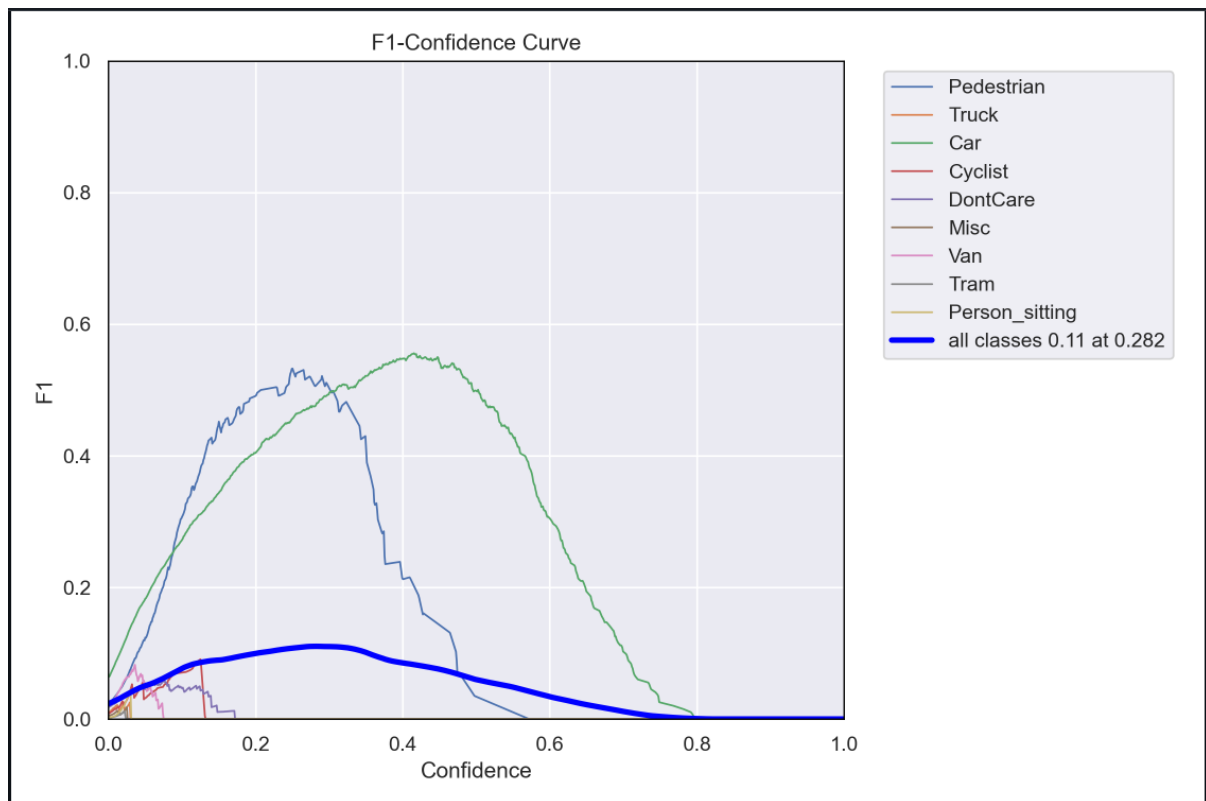
```
path: D:\present\cmpe249-hw1\dataset
train: D:\present\cmpe249-hw1\dataset\images\train
val: D:\present\cmpe249-hw1\dataset\images\test
test:
```

```
# Classes
names:
  0: Pedestrian
  1: Truck
  2: Car
  3: Cyclist
  4: DontCare
  5: Misc
  6: Van
  7: Tram
  8: Person_sitting
```

As the training took a long time for each epoch, ran the model for only 3 epochs with the command:
python train.py --epochs 3 --data coco.yaml --weights yolov5s.pt

The results can be seen [here](#).

For 3 epochs:



Step 4: Evaluate coco metrics

To evaluate coco metrics, run the [validation script, val.py](#) on the generated model, [best.pt](#)
`python val.py --weights best.pt --data coco.yaml`

Changes made in val.py,

```
save_json = True and
anno_json = '../dataset/coco-annotations/annotations_test.json'
```

Note: annotations_test.json is created as described in Step 1 of Inference – Fast RCNN section.
mAP50 values:

Class	Images	Instances	P	R	mAP50	mAP50-95:
all	100	671	0.892	0.122	0.116	0.0394
Pedestrian	100	55	0.643	0.427	0.399	0.11
Truck	100	14	1	0	0.00745	0.00218
Car	100	387	0.383	0.674	0.541	0.208
Cyclist	100	18	1	0	0.0236	0.0107
DontCare	100	139	1	0	0.0158	0.00451
Misc	100	9	1	0	0.00356	0.000927
Van	100	45	1	0	0.0297	0.00884
Tram	100	3	1	0	0.00251	0.0015
Person sitting	100	1	1	0	0.0255	0.00765

INFERENCE – FASTERRCNN

Step 1: Convert the KITTI dataset to COCO format

Fast RCNN requires COCO annotations for inference. Following these references, <https://medium.com/codable/convert-any-dataset-to-coco-object-detection-format-with-sahi-95349e1fe2b7> and <https://pypi.org/project/sahi/> converted the KITTI dataset to COCO format. The following is the summary of the code.

- Encode the categories from the labels files.
- For every training label, create a COCOImage and for every annotation in that label, create a COCOAnnotation object with the bounding boxes information (mid point and height and width of the box) and the category.
- Similarly, perform the same for test labels.
- Save both in .json format.

They can be found in the dataset directory [here](#).

dataset

-coco-annotations

- annotations_test.json

- annotations_train.json

Step 2: [Perform Inference \(Code\)](#)

Next imported a pre-trained Fast RCNN model from torchvision.models.detection and performed inference and generated a predictions.json in COCO format.

Note: The following transformations need to be done on predictions to stay consistent with the process in Step 1.

```
x1 = float(box[0])
```

```
y1 = float(box[1])
```

```
x2 = float(box[2])
```

```
y2 = float(box[3])
```

```
intx1 = int(x1)
```

```
inty1 = int(y1)
```

```
intx2 = int(x2)
```

```
inty2 = int(y2)
```

```
bbox_center_x = float( (x1 + (x2 - x1) / 2.0) / img_width)
```

```

bbox_center_y = float( (y1 + (y2 - y1) / 2.0) / img_height)
bbox_width = float((x2 - x1) / img_width)
bbox_height = float((y2 - y1) / img_height)

```

Step 3: [COCO Metrics Evaluation](#)

Used pycocotools for calculating COCO Evaluation metrics as seen from this reference:

<https://github.com/cocodataset/cocoapi/blob/master/PythonAPI/pycocoEvalDemo.ipynb>

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.026
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.074
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.000
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= 0.026
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.028
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.029
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.029
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= 0.029
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= -1.000

INFERENCE – YOLOv5

YOLOv5 follows a similar pattern to FasterRCNN inference.

- Generate the predictions json file on the test images using pretrained YOLOv5 model.

[Code](#) and [results](#)

Reference: https://pytorch.org/hub/ultralytics_yolov5/

- [Calculate COCO metrics.](#)

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.023
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.065
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.008
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= 0.023
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.019
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.035
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.035
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= 0.035
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= -1.000

INFERENCE PIPELINE

Code: <https://github.com/uttej कुमारreddy/cmpe249-hw1/blob/master/inference-pipeline/infer.py>

Running the inference pipeline for 2 models, [pretrained-yolov5](#) and [custom-trained-yolov5](#).

Model	Image	Output Latency	Result
Custom	001000.png	0:00:11.17	 A street scene with buildings on both sides. Three cars are detected with orange bounding boxes and labels: 'Car 0.46', 'Car 0.55', and 'Car 0.45'.
Custom	001001.png	0:00:11.11	 A street scene with a stone wall on the right and trees on the left. One car is detected with an orange bounding box and label: 'Car 0.58'.
Custom	001002.png	0:00:11.22	 A street scene with buildings on the left and trees on the right. Three cars are detected with orange bounding boxes and labels: 'Car 0.57', 'Car 0.51', and 'Car 0.36'.
Pretrained	001000.png	0:00:10.89	 A street scene with buildings on both sides. Three cars are detected with orange bounding boxes and labels: 'Car 0.78', 'Car 0.75', and 'Car 0.91'. Two traffic lights are detected with green bounding boxes and labels: 'traffic light 0.47' and 'traffic light 0.40'.
Pretrained	001001.png	0:00:11.28	 A street scene with a stone wall on the right and trees on the left. One car is detected with an orange bounding box and label: 'Car 0.76'.
Pretrained	001002.png	0:00:10.85	 A street scene with buildings on the left and trees on the right. Five cars are detected with orange bounding boxes and labels: 'Car 0.85', 'Car 0.71', 'Car 0.65', 'Car 0.42', and 'Car 0.75'.

References:

- [1] <https://medium.com/codable/convert-any-dataset-to-coco-object-detection-format-with-sahi-95349e1fe2b7>
- [2] <https://pypi.org/project/sahi/>
- [3] <https://github.com/packyan/Kitti2Coco/blob/master/kitti2coco-label-trans.py>
- [4] <https://github.com/cocodataset/cocoapi/blob/master/PythonAPI/pycocoEvalDemo.ipynb>
- [5] https://pytorch.org/hub/ultralytics_yolov5/