

CMPE 249

Report for HW1 – 2D Object Detection

Uttej Kumar Reddy Gade
016065543

Option: Training (1 model) & Inference (2 models)

Models for Training:

YOLOv5 (<https://github.com/ultralytics/yolov5>)

Models for Inference:

FastRCNN (from torchvision.models.detection import fasterrcnn_resnet50_fpn)

YOLOv5 (torch.hub.load('ultralytics/yolov5', 'yolov5s', pretrained=True))

Dataset: [KITTI 2D Object Detection Dataset](#)

Description:

Step 1: Reducing KITTI Dataset for quick training and inference

Extracted the first 1000 images and labels for training and the next 100 for testing from below links.

[Download left color images of object data set \(12 GB\)](#)

[Download training labels of object data set \(5 MB\)](#)

Organised the files in a manner suitable for YOLOv5.

Directory structure

```
dataset
- images
  -train
  -test
- labels-raw
  -train
  -test
```

TRAINING – YOLOv5

Step 2: Formatting dataset into suitable format for YOLOv5 training

YOLOv5 looks for labels in the format [CATEGORY] [BBOX_X] [BBOX_Y] [HEIGHT] [WIDTH] where BBOX_X and BBOX_Y are the center coordinates of the bounding box. However, the labels downloaded from the source has [CATEGORY] with 14 values encoding more attributes.

So, as the next step, converted the labels to required format using the below link as reference:

<https://github.com/packyan/Kitti2Coco/blob/master/kitti2coco-label-trans.py>

Code:

<https://github.com/uttej कुमारreddy/cmpe249-hw1/blob/master/kitti-labels-to-coco-format.ipynb>

Summary:

Encode all categories in the label to a number and store it.

- a. For each line in label file, extract the 5th-8th value and scale them to get required values.

```
bbox_center_x = float( (x1 + (x2 - x1) / 2.0) / img_width)
bbox_center_y = float( (y1 + (y2 - y1) / 2.0) / img_height)
bbox_width = float((x2 - x1) / img_width)
bbox_height = float((y2 - y1) / img_height)
```

- b. Write them in new label files in the above mentioned format.

Raw Labels:

Pedestrian 0.00 0 -0.20 712.40 143.00 810.73 307.92 1.89 0.48 1.20 1.84 1.47 8.41 0.01

Processed Labels:

0 0.6221936274509804 0.6093513513513513 0.08033496732026148 0.4457297297297298

The processed labels are stored [here](#).

dataset
- labels

Step 3: Downloaded, configured and ran YOLOv5

[Downloaded YOLOv5](#) from Github. Configured the coco.yaml file in data folder as follows:

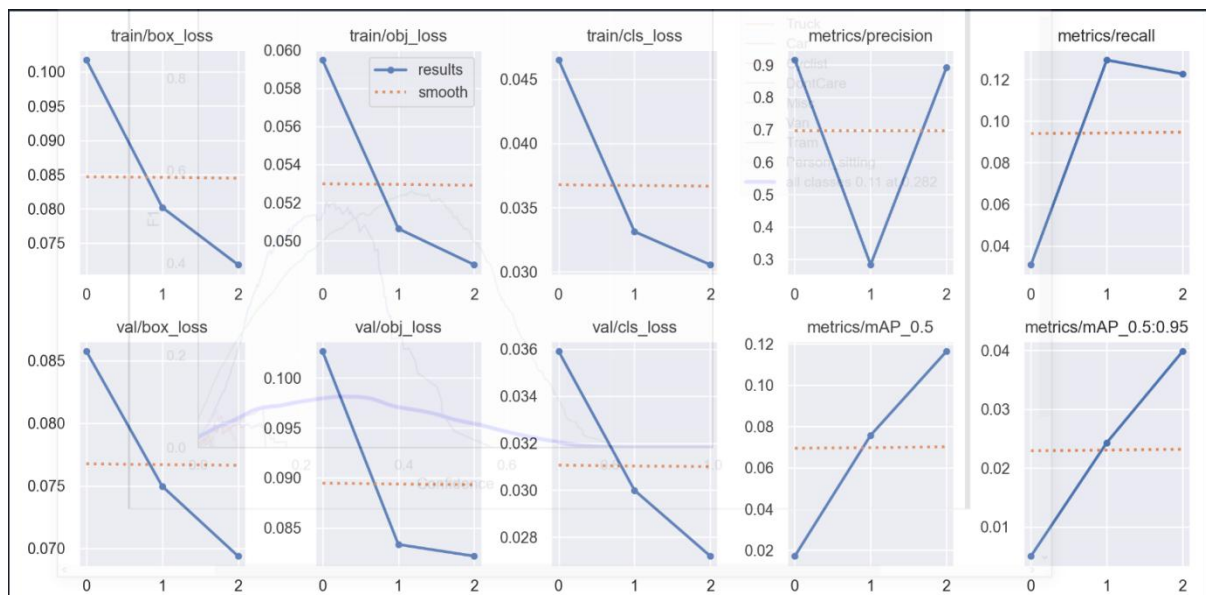

```
path: D:\present\cmpe249-hw1\dataset
train: D:\present\cmpe249-hw1\dataset\images\train
val: D:\present\cmpe249-hw1\dataset\images\test
test:
```

```
# Classes
names:
  0: Pedestrian
  1: Truck
  2: Car
  3: Cyclist
  4: DontCare
  5: Misc
  6: Van
  7: Tram
  8: Person_sitting
```

As the training took a long time for each epoch, ran the model for only 3 epochs with the command:
python train.py --epochs 3 --data coco.yaml --weights yolov5s.pt

The results can be seen [here](#).

For 3 epochs:



Step 4: Evaluate coco metrics

To evaluate coco metrics, run the [validation script, val.py](#) on the generated model, [best.pt](#)
`python val.py --weights best.pt --data coco.yaml`

Changes made in val.py,

```
save_json = True and
anno_json = '../dataset/coco-annotations/annotations_test.json'
```

Note: annotations_test.json is created as described in Step 1 of Inference – Fast RCNN section.
mAP50 values:

| Class | Images | Instances | P | R | mAP50 | mAP50-95: |
|----------------|--------|-----------|-------|-------|---------|-----------|
| all | 100 | 671 | 0.892 | 0.122 | 0.116 | 0.0394 |
| Pedestrian | 100 | 55 | 0.643 | 0.427 | 0.399 | 0.11 |
| Truck | 100 | 14 | 1 | 0 | 0.00745 | 0.00218 |
| Car | 100 | 387 | 0.383 | 0.674 | 0.541 | 0.208 |
| Cyclist | 100 | 18 | 1 | 0 | 0.0236 | 0.0107 |
| DontCare | 100 | 139 | 1 | 0 | 0.0158 | 0.00451 |
| Misc | 100 | 9 | 1 | 0 | 0.00356 | 0.000927 |
| Van | 100 | 45 | 1 | 0 | 0.0297 | 0.00884 |
| Tram | 100 | 3 | 1 | 0 | 0.00251 | 0.0015 |
| Person sitting | 100 | 1 | 1 | 0 | 0.0255 | 0.00765 |

INFERENCE – FASTERRCNN

Step 1: Convert the KITTI dataset to COCO format

Fast RCNN requires COCO annotations for inference. Following these references, <https://medium.com/codable/convert-any-dataset-to-coco-object-detection-format-with-sahi-95349e1fe2b7> and <https://pypi.org/project/sahi/> converted the KITTI dataset to COCO format. The following is the summary of the code.

- Encode the categories from the labels files.
- For every training label, create a COCOImage and for every annotation in that label, create a COCOAnnotation object with the bounding boxes information (mid point and height and width of the box) and the category.
- Similarly, perform the same for test labels.
- Save both in .json format.

They can be found in the dataset directory [here](#).

dataset

-coco-annotations

- annotations_test.json

- annotations_train.json

Step 2: [Perform Inference \(Code\)](#)

Next imported a pre-trained Fast RCNN model from torchvision.models.detection and performed inference and generated a predictions.json in COCO format.

Note: The following transformations need to be done on predictions to stay consistent with the process in Step 1.

```
x1 = float(box[0])
```

```
y1 = float(box[1])
```

```
x2 = float(box[2])
```

```
y2 = float(box[3])
```

```
intx1 = int(x1)
```

```
inty1 = int(y1)
```

```
intx2 = int(x2)
```

```
inty2 = int(y2)
```

```
bbox_center_x = float( (x1 + (x2 - x1) / 2.0) / img_width)
```

```

bbox_center_y = float( (y1 + (y2 - y1) / 2.0) / img_height)
bbox_width = float((x2 - x1) / img_width)
bbox_height = float((y2 - y1) / img_height)

```

Step 3: [COCO Metrics Evaluation](#)

Used pycocotools for calculating COCO Evaluation metrics as seen from this reference:

<https://github.com/cocodataset/cocoapi/blob/master/PythonAPI/pycocoEvalDemo.ipynb>

| | | |
|-------------------|---|----------|
| Average Precision | (AP) @[IoU=0.50:0.95 area= all maxDets=100] | = 0.026 |
| Average Precision | (AP) @[IoU=0.50 area= all maxDets=100] | = 0.074 |
| Average Precision | (AP) @[IoU=0.75 area= all maxDets=100] | = 0.000 |
| Average Precision | (AP) @[IoU=0.50:0.95 area= small maxDets=100] | = 0.026 |
| Average Precision | (AP) @[IoU=0.50:0.95 area=medium maxDets=100] | = -1.000 |
| Average Precision | (AP) @[IoU=0.50:0.95 area= large maxDets=100] | = -1.000 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets= 1] | = 0.028 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets= 10] | = 0.029 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets=100] | = 0.029 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= small maxDets=100] | = 0.029 |
| Average Recall | (AR) @[IoU=0.50:0.95 area=medium maxDets=100] | = -1.000 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= large maxDets=100] | = -1.000 |

INFERENCE – YOLOv5

YOLOv5 follows a similar pattern to FasterRCNN inference.

- Generate the predictions json file on the test images using pretrained YOLOv5 model.

[Code](#) and [results](#)

Reference: https://pytorch.org/hub/ultralytics_yolov5/

- [Calculate COCO metrics.](#)

| | | |
|-------------------|---|----------|
| Average Precision | (AP) @[IoU=0.50:0.95 area= all maxDets=100] | = 0.023 |
| Average Precision | (AP) @[IoU=0.50 area= all maxDets=100] | = 0.065 |
| Average Precision | (AP) @[IoU=0.75 area= all maxDets=100] | = 0.008 |
| Average Precision | (AP) @[IoU=0.50:0.95 area= small maxDets=100] | = 0.023 |
| Average Precision | (AP) @[IoU=0.50:0.95 area=medium maxDets=100] | = -1.000 |
| Average Precision | (AP) @[IoU=0.50:0.95 area= large maxDets=100] | = -1.000 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets= 1] | = 0.019 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets= 10] | = 0.035 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= all maxDets=100] | = 0.035 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= small maxDets=100] | = 0.035 |
| Average Recall | (AR) @[IoU=0.50:0.95 area=medium maxDets=100] | = -1.000 |
| Average Recall | (AR) @[IoU=0.50:0.95 area= large maxDets=100] | = -1.000 |

INFERENCE PIPELINE

Code: <https://github.com/uttej कुमारreddy/cmpe249-hw1/blob/master/inference-pipeline/infer.py>

Summary:

- a. Created a class InferencePipeline which initializes with model type (yolo-pretrained, yolo-custom, rcnn) and image path. The inputs are passed as args from command line.
- b. The inferencepipeline follows the steps:
 - a. Prechecks: Check if model and image inputs are valid
 - b. Load models: Load the yolo-pretrained, yolo-custom or rcnn model if they are valid.
 - c. Transform image: RCNN requires the image to be transformed to a tensor for prediction. If the model type if YOLO, this function does nothing.
 - d. Predict: Model is run on the transformed image.
 - e. Visualize: The images with bounding boxes, labels and scores are generated.

Running the inference pipeline for 2 models, [pretrained-yolov5](#), [custom-trained-yolov5](#) and fasterrcnn.

| Model | Image | Output Latency (s) | Result |
|-------------------|------------|--------------------|--|
| Custom Yolov5 | 001000.png | 6.98 |  A street scene with several cars. Bounding boxes and scores are shown: 'Car 0.46', 'Car 0.55', and 'Car 0.45'. |
| Custom Yolov5 | 001001.png | 6.098 |  A road scene with a car in the distance. Bounding box and score are shown: 'Car 0.58'. |
| Custom Yolov5 | 001002.png | 6.02 |  A street scene with multiple cars. Bounding boxes and scores are shown: 'car 0.57', 'car 0.51', 'car 0.49', and 'car 0.38'. |
| Pretrained Yolov5 | 001000.png | 6.08 |  A street scene with cars and traffic lights. Bounding boxes and scores are shown: 'traffic light 0.47', 'car 0.78', 'car 0.75', 'car 0.91', and 'traffic light 0.40'. |

| | | | |
|-------------------|------------|------|--|
| Pretrained YOLOv5 | 001001.png | 6.28 |  |
| Pretrained YOLOv5 | 001002.png | 5.93 |  |
| FasterRCNN | 001000.png | 2.21 | |
| FasterRCNN | 001001.png | 2.18 | |
| FasterRCNN | 001002.png | 2.36 | |

The approximate average output latencies for the 3 models on the constructed inference pipeline and from the chosen three images is:

- a. Pretrained YOLO : 6.1 sec
- b. Custom YOLO : 6.37 sec
- c. FasterRCNN : 2.25 sec

References:

- [1] <https://medium.com/codable/convert-any-dataset-to-coco-object-detection-format-with-sahi-95349e1fe2b7>
- [2] <https://pypi.org/project/sahi/>
- [3] <https://github.com/packyan/Kitti2Coco/blob/master/kitti2coco-label-trans.py>
- [4] <https://github.com/cocodataset/cocoapi/blob/master/PythonAPI/pycocoEvalDemo.ipynb>
- [5] https://pytorch.org/hub/ultralytics_yolov5/